

Web Server Core Design (Part 1)



BY EVAN OLDS
CPT S 422

Needed Elements



- Proper concurrency model
- Ability to have a reusable request parser that can support different types of handlers
 - We don't want to have to re-write the core web server stuff when we make new web “applications”
 - Most web applications are not actually standalone apps. Rather, they are almost like plugins that get request objects from a core server
 - ✦ IIS
 - ✦ Apache

Motivation for Concurrency



- Web server needs to handle multiple requests at once
- Can't wait to complete download of 1 GB file before letting any other clients connect
- Server WILL only listen on 1 port, so in a sense it does accept one client at a time
 - More of a literal sense here: at any given moment we're blocked on a single [AcceptTcpClient](#) call in the “listening thread”
- What logical question does this lead you to ask about [AcceptTcpClient](#)?

What AcceptTcpClient Gives Us



- “After accepting a TcpClient with AcceptTcpClient can I send the client object off to some other thread and go right back to another AcceptTcpClient call on the listening thread?”
- Yes. AcceptTcpClient, if it succeeds, gives us a socket that we can communicate with while accepting another client on the same port as the previous AcceptTcpClient call.
 - The TcpClient returned basically communicates on a different port after it has been accepted, so the original remains open

Process requests on new...?



- So we can read from and write to a TcpClient while waiting to accept another
- This means we can process the request on a new thread (if one is available)
 - Threads from thread pool should be waiting, ready to handle requests from accepted TCP clients
- We WILL use threads and a thread pool for our server, but this isn't necessarily what's used on the average industry web server
- What the other option for a concurrency mechanism that could process a request?

Threads vs. Processes



- Many web servers will spawn a new process to deal with a request
 - Why?
 - What does using a process over a thread allow us to do?
 - So that we see both pros and cons, what would be the advantage of threads only instead of processes?
 - Discuss in class

Threads vs. Processes



- **Reasons threads are better**
 - Faster. Thread pool gives us quick access to threads.
 - Easier to manage a thread pool vs. a process pool. Don't need IPC. Also, the process-based-handling servers actually are not likely to use a process pool because of reasons that should have been discussed from the previous slide (need to spawn process with an account that's related to the permissions associated with the request)
- **Reason processes are better**
 - Process permissions allow for use of operating system security. Restricts the request from accessing areas of the file system that it should have access to (i.e. "sandboxed")
 - If process crashes, not as big of a deal as a thread crashing.

Listening Thread Logic



```
while (server_active) {  
    1. Accept TCP client  
    2. Send client to a thread pool thread to be processed  
}
```

- More to come in this logic when we start discussing how to terminate the server
 - Raises the question of how we break out of blocking accept call

Worker Thread Logic



```
while (server_active) {  
    TcpClient client = sharedBlockingCollection.Take();  
    // Read request  
    // Write response (if applicable)  
    // Close client  
}
```

- Is it that simple?
- Questions we haven't addressed yet:
 - What do we write in the response?
 - How do we design a reusable server code that allows for different types of handlers to be written?

Worker Thread Logic



- Consider:

```
while (server_active) {  
    TcpClient client = sharedBlockingCollection.Take();  
    MyWebRequest req = PackageRequest(client);  
    MyWebHandler handler = FindHandlerFor(req);  
    handler.Handle(req);  
    // Close client  
}
```

Parsing and Packaging the Request



- Want the ability to have different types of handlers
- Don't want each handler to have to do the following:
 - Parse HTTP request to find out whether or not it's even valid
 - Parse HTTP request pieces to get the method, URI and version as separate strings
 - Parse HTTP headers and put them in some easily accessible collection
 - Put the body of the request into a simple Stream object
- Core logic, in a “PackageRequest” or “BuildRequest” function does all the above.
 - Send “simple” (as it can be) request object to handler
 - Also implies that the core logic manages a collection of handlers and determines which ones handle which requests

Request Object



- Request object that we pass to handlers contains:
 - Parsed info from the first line of the request
 - ✦ Method (GET, HEAD, POST, PUT, DELETE, TRACE, OPTIONS, CONNECT, or PATCH)
 - ✦ Request target (URI)
 - ✦ HTTP version
 - Headers parsed and packaged (what data structure is good to store a collection of header names/values?)
 - Stream for the body
- Request object may also provide utility functions to write responses

Request Object



```
public class Request
{
    Stream Body;
    ConcurrentDictionary<string, string> Headers;
    string Method;
    string RequestTarget;
    string HTTPVersion;
    public long GetContentLengthOrDefault(long defaultValue)
    {
        // ??
    }
    public Tuple<long, long> GetRangeHeader()
    {
        // ??
    }
    System.Net.Sockets.NetworkStream Response;
}
```

Request Object Body Stream



- There is content in an HTTP request before the body (request line and headers)
- Want a single Stream object for the body
 - This stream doesn't contain any data before the body
 - First byte in stream is first byte of body content
 - Wouldn't be convenient if it were any other way
- If we're implementing a BuildRequest function that needs to set the body stream, where do we get it from?
 - Can't just set it to the network stream. Why not?

Request Object Body Stream



- If the position in the network stream were right at the beginning of the body, then we COULD just set it as the body stream
 - Shows us that the body stream is likely not going to support seeking
 - Forward-only reading
- What if we read into a buffer and got the entire first line, all the headers, and some of the body?
- How do we construct a Stream that includes the part of the data we've already read as well as content from the NetworkStream that has yet to be read?